

```
import os
import cv2
import numpy as np
```

Reading Input Images

```
images = []
labels = []
trial = []
```

Image Augmentation

Daisy Inputs

```
for img in os.listdir('/home/wintermute/Downloads/flowers/daisy/') :
    label = 0
    path = os.path.join('/home/wintermute/Downloads/flowers/daisy/',
img)
    img = cv2.imread(path, cv2.IMREAD_COLOR)
    img = cv2.resize(img, (150, 150))
    aug_img = cv2.flip(img,1)
    images.append(np.array(img))
    images.append(np.array(aug_img))
    labels.append(str(label))
    labels.append(str(label))
```

Dandelion Inputs

```
for img in os.listdir('/home/wintermute/Downloads/flowers/dandelion/') :
    label = 1
    path =
os.path.join('/home/wintermute/Downloads/flowers/dandelion/', img)
    img = cv2.imread(path, cv2.IMREAD_COLOR)
    img = cv2.resize(img, (150, 150))
    aug_img = cv2.flip(img,1)
    images.append(np.array(img))
    images.append(np.array(aug_img))
    labels.append(str(label))
    labels.append(str(label))
```

Rose Inputs

```
for img in os.listdir('/home/wintermute/Downloads/flowers/rose/') :
    label = 2
    path = os.path.join('/home/wintermute/Downloads/flowers/rose/',
img)
    img = cv2.imread(path, cv2.IMREAD_COLOR)
    img = cv2.resize(img, (150, 150))
    aug_img = cv2.flip(img,1)
    images.append(np.array(img))
    images.append(np.array(aug_img))
```

```
labels.append(str(label))
labels.append(str(label))
```

Sunflower Inputs

```
for img in os.listdir('/home/wintermute/Downloads/flowers/sunflower/'):
    label = 3
    path =
os.path.join('/home/wintermute/Downloads/flowers/sunflower/', img)
img = cv2.imread(path, cv2.IMREAD_COLOR)
img = cv2.resize(img, (150, 150))
aug_img = cv2.flip(img,1)
images.append(np.array(img))
images.append(np.array(aug_img))
labels.append(str(label))
labels.append(str(label))
```

Tulip Inputs

```
for img in os.listdir('/home/wintermute/Downloads/flowers/tulip/'):
    label = 4
    path = os.path.join('/home/wintermute/Downloads/flowers/tulip/',
img)
img = cv2.imread(path, cv2.IMREAD_COLOR)
img = cv2.resize(img, (150, 150))
aug_img = cv2.flip(img,1)
images.append(np.array(img))
images.append(np.array(aug_img))
labels.append(str(label))
labels.append(str(label))
```

```
print(len(images))
print(len(labels))
```


```
8634
8634
```

```
images = np.array(images)
```

Visualizing Image

```
import matplotlib.pyplot as plt

plt.imshow(images[6])
plt.show()
```



Data Preparation

```
from sklearn.model_selection import *  
import keras
```

```
2022-10-03 12:57:10.016905: I  
tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow  
binary is optimized with oneAPI Deep Neural Network Library (oneDNN)  
to use the following CPU instructions in performance-critical  
operations: AVX2 FMA  
To enable them in other operations, rebuild TensorFlow with the  
appropriate compiler flags.  
2022-10-03 12:57:11.122756: E  
tensorflow/stream_executor/cuda/cuda_blas.cc:2981] Unable to register  
cuBLAS factory: Attempting to register factory for plugin cuBLAS when  
one has already been registered  
2022-10-03 12:57:14.551792: W  
tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could  
not load dynamic library 'libnvinfer.so.7'; dLError: libnvinfer.so.7:  
cannot open shared object file: No such file or directory;  
LD_LIBRARY_PATH: /home/wintermute/.local/lib/python3.10/site-  
packages/cv2/../../lib64:  
2022-10-03 12:57:14.552228: W  
tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could  
not load dynamic library 'libnvinfer_plugin.so.7'; dLError:
```

```
libnvinfer_plugin.so.7: cannot open shared object file: No such file
or directory; LD_LIBRARY_PATH:
/home/wintermute/.local/lib/python3.10/site-packages/cv2/../../lib64:
2022-10-03 12:57:14.552245: W
tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning:
Cannot dlopen some TensorRT libraries. If you would like to use Nvidia
GPU with TensorRT, please make sure the missing libraries mentioned
above are installed properly.
```

```
labels = keras.utils.to_categorical(labels,num_classes = 5)
```

```
print(labels)
```

```
[[1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0.]
 ...
 [0. 0. 0. 0. 1.]
 [0. 0. 0. 0. 1.]
 [0. 0. 0. 0. 1.]]
```

```
x_train,x_test,y_train,y_test =
train_test_split(images,labels,test_size = 0.2,random_state = 42)
```

```
x_train.shape
```

```
(6907, 150, 150, 3)
```

```
x_train,x_val,y_train,y_val =
train_test_split(x_train,y_train,test_size = 0.15,random_state = 42)
```

```
import keras
from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout, Flatten
from keras.layers import Conv2D
from keras.layers import MaxPooling2D,MaxPool2D
from keras.optimizers import Adam
from tensorflow.keras.layers import Normalization
from tensorflow.keras.layers import LeakyReLU
from tensorflow.keras.layers import Conv2D, MaxPooling2D,
GlobalMaxPooling2D
```

```
model = Sequential()
```

```
model.add(Conv2D(128, (3, 3), input_shape=x_train.shape[1:]))
model.add(LeakyReLU(alpha=0.02))
```

```
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
```

```
model.add(Conv2D(128, (3, 3)))
model.add(LeakyReLU(alpha=0.02))
```

```
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
```

```
model.add(GlobalMaxPooling2D())
```

```
model.add(Dense(512))
model.add(LeakyReLU(alpha=0.02))
model.add(Dropout(0.5))
```

```
model.add(Dense(5))
model.add(Activation('softmax'))
```

```
2022-10-03 12:57:30.868525: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful
NUMA node read from SysFS had negative value (-1), but there must be
at least one NUMA node, so returning NUMA node zero
```

```
2022-10-03 12:57:30.998271: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful
NUMA node read from SysFS had negative value (-1), but there must be
at least one NUMA node, so returning NUMA node zero
```

```
2022-10-03 12:57:30.999508: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful
NUMA node read from SysFS had negative value (-1), but there must be
at least one NUMA node, so returning NUMA node zero
```

```
2022-10-03 12:57:31.003081: I
tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow
binary is optimized with oneAPI Deep Neural Network Library (oneDNN)
to use the following CPU instructions in performance-critical
operations: AVX2 FMA
```

To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

```
2022-10-03 12:57:31.012538: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful
NUMA node read from SysFS had negative value (-1), but there must be
at least one NUMA node, so returning NUMA node zero
```

```
2022-10-03 12:57:31.013477: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful
NUMA node read from SysFS had negative value (-1), but there must be
at least one NUMA node, so returning NUMA node zero
```

```
2022-10-03 12:57:31.013641: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful
NUMA node read from SysFS had negative value (-1), but there must be
at least one NUMA node, so returning NUMA node zero
```

```
2022-10-03 12:57:33.469275: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful
NUMA node read from SysFS had negative value (-1), but there must be
at least one NUMA node, so returning NUMA node zero
```

```
2022-10-03 12:57:33.470573: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful
```

NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero

2022-10-03 12:57:33.472435: I

tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:980] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero

2022-10-03 12:57:33.473338: I

tensorflow/core/common_runtime/gpu/gpu_device.cc:1616] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 2121 MB memory: -> device: 0, name: NVIDIA GeForce RTX 3050 Laptop GPU, pci bus id: 0000:02:00.0, compute capability: 8.6

```
model.compile(loss='categorical_crossentropy', optimizer='adam',  
metrics=['accuracy'])
```

```
history = model.fit(x_train, y_train, epochs=25, batch_size=32,  
validation_data=(x_val, y_val))
```

2022-10-03 12:57:42.446183: W

tensorflow/core/framework/cpu_allocator_impl.cc:82] Allocation of 396225000 exceeds 10% of free system memory.

2022-10-03 12:57:43.224683: W

tensorflow/core/framework/cpu_allocator_impl.cc:82] Allocation of 396225000 exceeds 10% of free system memory.

Epoch 1/25

2022-10-03 12:57:52.783676: I

tensorflow/stream_executor/cuda/cuda_dnn.cc:384] Loaded cuDNN version 8500

2022-10-03 12:57:55.865772: I

tensorflow/stream_executor/cuda/cuda_blas.cc:1614] TensorFloat-32 will be used for the matrix multiplication. This will only be logged once.

184/184 [=====] - 53s 203ms/step - loss: 5.4625 - accuracy: 0.4342 - val_loss: 1.0550 - val_accuracy: 0.6316
Epoch 2/25

184/184 [=====] - 36s 197ms/step - loss: 1.0916 - accuracy: 0.5644 - val_loss: 0.9466 - val_accuracy: 0.6577

Epoch 3/25

184/184 [=====] - 41s 222ms/step - loss: 1.0087 - accuracy: 0.6060 - val_loss: 0.9776 - val_accuracy: 0.6528

Epoch 4/25

184/184 [=====] - 41s 222ms/step - loss: 0.9426 - accuracy: 0.6390 - val_loss: 0.8582 - val_accuracy: 0.6818

Epoch 5/25

184/184 [=====] - 44s 241ms/step - loss: 0.9008 - accuracy: 0.6508 - val_loss: 0.8867 - val_accuracy: 0.6606

Epoch 6/25

184/184 [=====] - 63s 342ms/step - loss: 0.8671 - accuracy: 0.6681 - val_loss: 0.8365 - val_accuracy: 0.6644

Epoch 7/25
184/184 [=====] - 82s 445ms/step - loss:
0.8294 - accuracy: 0.6794 - val_loss: 0.7286 - val_accuracy: 0.7406
Epoch 8/25
184/184 [=====] - 89s 485ms/step - loss:
0.8116 - accuracy: 0.6894 - val_loss: 0.7453 - val_accuracy: 0.7387
Epoch 9/25
184/184 [=====] - 89s 485ms/step - loss:
0.8066 - accuracy: 0.6918 - val_loss: 0.7996 - val_accuracy: 0.7040
Epoch 10/25
184/184 [=====] - 89s 485ms/step - loss:
0.7702 - accuracy: 0.7089 - val_loss: 0.7026 - val_accuracy: 0.7387
Epoch 11/25
184/184 [=====] - 89s 483ms/step - loss:
0.7501 - accuracy: 0.7121 - val_loss: 0.7037 - val_accuracy: 0.7512
Epoch 12/25
184/184 [=====] - 64s 345ms/step - loss:
0.7370 - accuracy: 0.7208 - val_loss: 0.7229 - val_accuracy: 0.7367
Epoch 13/25
184/184 [=====] - 45s 246ms/step - loss:
0.7489 - accuracy: 0.7155 - val_loss: 0.7112 - val_accuracy: 0.7358
Epoch 14/25
184/184 [=====] - 43s 235ms/step - loss:
0.7067 - accuracy: 0.7358 - val_loss: 0.7007 - val_accuracy: 0.7338
Epoch 15/25
184/184 [=====] - 47s 256ms/step - loss:
0.7255 - accuracy: 0.7249 - val_loss: 0.7051 - val_accuracy: 0.7464
Epoch 16/25
184/184 [=====] - 53s 287ms/step - loss:
0.7180 - accuracy: 0.7242 - val_loss: 0.5987 - val_accuracy: 0.7772
Epoch 17/25
184/184 [=====] - 89s 484ms/step - loss:
0.6625 - accuracy: 0.7450 - val_loss: 0.6584 - val_accuracy: 0.7551
Epoch 18/25
184/184 [=====] - 89s 485ms/step - loss:
0.6918 - accuracy: 0.7378 - val_loss: 0.6580 - val_accuracy: 0.7647
Epoch 19/25
184/184 [=====] - 89s 484ms/step - loss:
0.6697 - accuracy: 0.7514 - val_loss: 0.6395 - val_accuracy: 0.7512
Epoch 20/25
184/184 [=====] - 89s 484ms/step - loss:
0.6560 - accuracy: 0.7537 - val_loss: 0.7107 - val_accuracy: 0.7358
Epoch 21/25
184/184 [=====] - 89s 484ms/step - loss:
0.6440 - accuracy: 0.7537 - val_loss: 0.7112 - val_accuracy: 0.7107
Epoch 22/25
184/184 [=====] - 89s 484ms/step - loss:
0.6325 - accuracy: 0.7549 - val_loss: 0.6934 - val_accuracy: 0.7329
Epoch 23/25
184/184 [=====] - 89s 485ms/step - loss:

0.6503 - accuracy: 0.7506 - val_loss: 0.6827 - val_accuracy: 0.7531
Epoch 24/25

184/184 [=====] - 89s 484ms/step - loss:
0.6375 - accuracy: 0.7629 - val_loss: 0.6367 - val_accuracy: 0.7541
Epoch 25/25

184/184 [=====] - 89s 484ms/step - loss:
0.6524 - accuracy: 0.7560 - val_loss: 0.6443 - val_accuracy: 0.7676


```
print("Test Accuracy: {0:.2f}%".format(model.evaluate(x_test,y_test)
[1]*100))
```

54/54 [=====] - 6s 103ms/step - loss: 0.6752
- accuracy: 0.7678

Test Accuracy: 76.78%

```
model.save_weights('Assignment_3_Final_Weights.h5')
```

```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epochs')
plt.legend(['train', 'test'])
plt.show()
```




```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epochs')
plt.legend(['train', 'test'])
plt.show()
```

